

COMPUTERS, ETHICS, AND SOCIETY

THIRD EDITION

Edited by

M. David Ermann

Michele S. Shauf

New York Oxford

OXFORD

UNIVERSITY PRESS

2003

Oxford University Press

Oxford New York

Auckland Bangkok Buenos Aires Cape Town Chennai
Dar es Salaam Delhi Hong Kong Istanbul Karachi Kolkata
Kuala Lumpur Madrid Melbourne Mexico City Mumbai
Nairobi São Paulo Shanghai Taipei Tokyo Toronto

Copyright © 1990, 1997, 2003 by Oxford University Press, Inc.

Published by Oxford University Press, Inc.

198 Madison Avenue, New York, New York, 10016

<http://www.oup-usa.org>

Oxford is a registered trademark of Oxford University Press

All rights reserved. No part of this publication may be reproduced,
stored in a retrieval system, or transmitted, in any form or by any means,
electronic, mechanical, photocopying, recording, or otherwise,
without the prior permission of Oxford University Press.

Library of Congress Cataloging-in-Publication Data

Computers, ethics, and society / edited by M. David Ermann, Michele S. Shauf.—3rd ed.
p. cm.

Includes bibliographical references and index.


ISBN 0-19-514302-7

1. Computers and civilization. 2. Computer security. 3. Human-computer interaction.
I. Ermann, M. David. II. Shauf, Michele S.

QA76.9.C66 C572 2003

303.48'34—dc21

2002072283



For Natalie, Mike, and Marlene

9 8 7 6 5 4 3 2 1

Printed in the United States of America
on acid-free paper

Philosophical Ethics

What constitutes an ethical choice? Does an ethical choice maximize happiness? Goodness? Does it follow some other universal principle? Is a universally applicable principle even possible? Ethical philosophers have long wrestled with these questions.

This section examines three ethical models: utilitarian, Kantian, and Nicomachean. In studying the following selections, consider how they might be applied to computer technologies. Following Aristotle's thinking, for example, can one conclude that hacking is ethical when it stems from a rational challenge conquered by a clever programmer? (See also Spafford, "Are Hacker Break-ins Ethical?")

1

The Best Action is the One with the Best Consequences

Of those actions available to you, you are morally obliged to choose that action which maximizes total happiness (summed over all affected persons) according to utilitarian ethical theory. The utilitarian model is particularly useful in illuminating instances when many people are affected in different ways by an action; for example, a utilitarian analysis may be useful in deciding what the laws ought to be on copyright (see National Research Council, "Music: Intellectual Property's Canary in the Digital Coal Mine") and privacy (see Garfinkel, "Privacy in a Database Nation").

John Hospers

Once one admits that one's own personal good is not the only consideration, how can one stop short of the good of everyone—"the general good"? This conclusion, at any rate, is the thesis of the ethical theory known as *utilitarianism*. The thesis is simply stated, though its application to actual situations is often extremely complex: whatever is intrinsically good should be promoted, and, accordingly, our obligation (or duty) is always to act so as to promote the greatest possible intrinsic good. It is never our duty to promote a lesser good when we could, by our action, promote a greater one; and the act which we should perform in any given situation is, therefore, the one which produces more intrinsic good than any other act we could have performed in its stead. In brief, the main tenet of utilitarianism is the maximization of intrinsic good.

Excerpt from *Human Conduct: Problems of Ethics*, Shorter Edition by John Hospers, Copyright © 1972 by Harcourt Brace & Company, reprinted by permission of the publisher.

The description just given is so brief that it will almost inevitably be misleading when one attempts to apply it in actual situations unless it is spelled out in greater detail. Let us proceed at once, then, to the necessary explanations and qualifications.

1. When utilitarians talk about right or wrong acts, they mean—and this point is shared by the proponents of all ethical theories—*voluntary acts*. Involuntary acts like the knee jerk are not included since we have no control over them: once the stimulus has occurred the act results quite irrespective of our own will. The most usual way in which the term “voluntary act” is defined is as follows:¹ an act is voluntary if the person *could* have acted differently *if* he had so chosen. For example, I went shopping yesterday, but if I had chosen (for one reason or another) to remain at home, I would have done so. My choosing made the difference. Making this condition is not the same as saying that an act, to be voluntary, must be *premeditated* or that it must be the outcome of *deliberation*, though voluntary acts often are planned. If you see a victim of a car accident lying in the street, you may rush to help him at once, without going through a process of deliberation; nevertheless your act is voluntary in that if you had chosen to ignore him you would have acted differently. Though not premeditated, the action was within your control. “Ought implies can,” and there is no ought when there is no can. To be right or wrong, an act must be within your power to perform: it must be performable as the result of your choice, and a different choice must have led to a different act or to no act at all.

2. There is no preference for immediate, as opposed to remote, happiness. If Act A will produce a certain amount of happiness today and Act B will produce twice as much one year hence, I should do B, even though its effects are more remote. Remoteness does not affect the principle at all: happiness is as intrinsically good tomorrow or next year as it is today, and one should forego a smaller total intrinsic good now in favor of a larger one in the future. (Of course, a remote happiness is often less certain to occur. But in that case we should choose A not because it is more immediate but because it is more nearly certain to occur.) . . .

3. Unhappiness must be considered as well as happiness. Suppose that Act A will produce five units of happiness and none of unhappiness and Act B will produce ten units of happiness and ten of unhappiness. Then A is to be preferred because the *net* happiness—the resulting total after the unhappiness has been subtracted from it—is greater in A than in B: it is five in A and zero in B. Thus the formula “You should do what will produce the greatest total happiness” is not quite accurate; you should do what will produce the most *net* happiness. This modification is what we shall henceforth mean in talking about “producing the greatest happiness”—we shall assume that the unhappiness has already been figured into the total.

4. It is not even accurate to say that you should always do what leads to

the greatest *balance* of happiness over unhappiness, for there may be no such balance in any alternative open to the agent: he may have to choose between “the lesser of two evils.” If Act A leads to five units of happiness and ten of unhappiness and Act B leads to five units of happiness and fifteen of unhappiness, you should choose A, not because it produces the most happiness (they both produce an equal amount) and not because there is a greater balance of happiness over unhappiness in A (there is a balance of unhappiness over happiness in both), but because, although both A and B produce a balance of unhappiness over happiness, A leads to a *smaller* balance of unhappiness over happiness than B does. Thus we should say, “Do that act which produces the greatest balance of happiness over unhappiness, or, if no act possible under the circumstance does this, do the one which produces the smallest balance of unhappiness over happiness.” This qualification also we shall assume to be included in the utilitarian formula from now on in speaking of “producing the greatest happiness” or “maximizing happiness.”

5. One should not assume that an act is right according to utilitarianism simply because it produces more happiness than unhappiness in its total consequences. If one did make this assumption, it would be right for ten men collectively to torture a victim, provided that the total pleasure enjoyed by the sadists exceeded the pain endured by the victim (assuming that pain is here equated with unhappiness and that all the persons died immediately thereafter and there were no further consequences). The requirement is not that the happiness exceed the unhappiness but that it do so *more* than any other act that could have been performed instead. This requirement is hardly fulfilled here: it is very probable indeed that the torturers could think of something better to do with their time.

6. When there is a choice between a greater happiness for yourself at the expense of others, and a greater happiness for others at the expense of your own, which should you choose? You choose, according to the utilitarian formula, whatever alternative results in the greater total amount of *net* happiness, precisely as we have described. If the net happiness is greater in the alternative favorable to yourself, you adopt this alternative; otherwise not. Mill says, “The happiness which forms the utilitarian standard of what is right in conduct, is not the agent’s own happiness, but that of all concerned. As between his own happiness and that of others, utilitarianism requires him to be as strictly impartial as a disinterested and benevolent spectator.”² To state this in different language, you are not to ignore your own happiness in your calculations, but neither are you to consider it more important than anyone else’s; you count as one, and only as one, along with everyone else. Thus if Act A produces a total net happiness of one hundred, and Act B produces seventy-five, A is the right act even if you personally would be happier in consequence of B. Your choice should not be an “interested”

one; you are not to be prejudiced in favor of your own happiness nor, for that matter, against it; your choice should be strictly *disinterested* as in the case of an impartial judge. Your choice should be dictated by the greatest-total-happiness principle, not by a *your-greatest-happiness* principle. If you imagine yourself as a judge having to make a decision designed to produce the most happiness for all concerned *without* knowing which of the people affected would be *you*, you have the best idea of the impartiality of judgment required by the utilitarian morality.

In egoistic ethics . . . your sole duty is to promote your own interests as much as possible, making quite sure, of course, that what you do will make you really happy (or whatever else you include in "your own interest") and that you do not choose merely what you *think* at the moment will do so; we have called this policy the policy of "*enlightened self-interest*." In an *altruistic* ethics, on the other hand, you sacrifice your own interests completely to those of others: you ignore your own welfare and become a doormat for the fulfillment of the interests of others. . . . But the utilitarian ethics is neither egoistic nor altruistic: it is a *universalistic* ethics, since it considers your interests equally with everyone else's. You are not the slave of others, nor are they your slaves. Indeed, there are countless instances in which the act required of you by ethical egoism and the act required by utilitarianism will be the same: for very often indeed the act that makes you happy will also make those around you happy, and by promoting your own welfare you will also be promoting theirs. (As support for this position, consider capitalistic society: the producer of wealth, by being free to amass profits, will have more incentive to produce and, by increasing production, will be able to create more work and more wealth. By increasing production, he will be increasing the welfare of his employees and the wealth of the nation.) Moreover, it is much more likely that you can effectively produce good by concentrating on your immediate environment than by "spreading yourself thin" and trying to help everyone in the world: "do-gooders" often succeed in achieving no good at all. (But, of course, sometimes they do.) You are in a much better position to produce good among those people whose needs and interests you already know than among strangers; and, of course, the person whose needs and interests you probably know best of all (though not always) is yourself. Utilitarianism is very far, then, from recommending that you ignore your own interests.

It is only when your interests cannot be achieved except at the cost of sacrificing the *greater* interests of others that utilitarianism recommends self-sacrifice. When interests conflict, you have to weigh your own interest against the general interest. If, on the one hand, you are spending all your valuable study time (and thus sacrificing your grades and perhaps your college degree) visiting your sick aunt because she wants you to, you would probably produce more good by spending your time studying. But on the

other hand, if an undeniably greater good will result from your sacrifice, if, for instance, your mother is seriously ill and no one else is available to care for her, you might have to drop out of school for a semester to care for her. It might even, on occasion, be your utilitarian duty to sacrifice your very life for a cause, when the cause is extremely worthy and requires your sacrifice for its fulfillment. But you must first make quite sure that your sacrifice will indeed produce the great good intended; otherwise you would be throwing your life away uselessly. You must act with your eyes open, not under the spell of a martyr complex.

.

7. The general temper of the utilitarian ethics can perhaps best be seen in its attitude toward moral rules, the traditional *dos* and *don'ts*. What is the utilitarian's attitude toward rules such as "Don't kill," "Don't tell lies," "Don't steal"?

According to utilitarianism, such rules are *on the whole* good, useful, and worthwhile, but they *may* have exceptions. None of them is sacrosanct. If killing is wrong, it is not because there is something intrinsically bad about killing itself, but because killing leads to a diminution of human happiness. This undesirable consequence almost always occurs: when a man takes another human life, he not only extinguishes in his victim all chances of future happiness, but he causes grief, bereavement, and perhaps years of misery for the victim's family and loved ones; moreover, for weeks or months countless people who know of this act may walk the streets in fear, wondering who will be the next victim—the amount of insecurity caused by even one act of murder is almost incalculable; and in addition to all this unhappiness, every violation of a law has a tendency to weaken the whole fabric of the law itself and tends to make other violations easier and more likely to occur. If the guilty man is caught, he himself hardly gains much happiness from lifelong imprisonment, nor are other people usually much happier for long because of his incarceration; and if he is not caught, many people will live in fear and dread, and he himself will probably repeat his act sooner or later, having escaped capture this time. The good consequences, if any, are few and far between and are overwhelmingly outweighed by the bad ones. Because of these prevailingly bad consequences, killing is condemned by the utilitarian, and thus he agrees with the traditional moral rule prohibiting it.

He would nevertheless admit the possibility of exceptions: if you had had the opportunity to assassinate Hitler in 1943 and did not, the utilitarian would probably say that you were doing wrong in *not* killing him. By not killing him, you would be stealing the death of thousands, if not millions, of other people: political prisoners and Jews whom he tortured and killed in concentration camps and thousands of soldiers (both Axis and Allied)

whose lives would have been saved by an earlier cessation of the war. If you had refrained from killing him when you had the chance, saying "It is my duty never to take a life, therefore I shall not take his," the man whose life you saved would then turn around and have a thousand others killed, and for his act the victims would have you to thank. Your conscience, guided by the traditional moral rules, would have helped to bring about the torture and death of countless other people.

Does the utilitarian's willingness to adopt violence upon occasion mean that a utilitarian could never be a pacifist? Not necessarily. He *might* say that *all* taking of human life is wrong, but if he took this stand, he would do so because he believed that *killing always leads to worse consequences (or greater unhappiness) than not killing and not because there is anything intrinsically bad about killing*. He might even be able to make out a plausible argument for saying that killing Hitler would have been wrong: perhaps even worse men would have taken over and the slaughter wouldn't have been prevented (but then wouldn't it have been right to kill *all* of them if one had the chance?); perhaps Hitler's "intuitions" led to an earlier defeat for Germany than if stabler men had made more rationally self-seeking decisions on behalf of Nazi Germany; perhaps the assassination of a bad leader would help lead to the assassination of a good one later on. With regard to some Latin American nations, at any rate, one might argue that killing one dictator would only lead to a revolution and another dictator just as bad as the first, with the consequent assassination of the second one, thus leading to revolution and social chaos and a third dictator. There are countless empirical facts that must be taken into consideration and carefully weighed before any such decision can safely be made. The utilitarian is not committed to saying that any one policy or line of action is the best in any particular situation, for what is best depends on empirical facts which may be extremely difficult to ascertain. All he is committed to is the statement that when the action is one that does not promote human happiness as much as another action that he could have performed instead, then the action is wrong; and that when it does promote more happiness, it is right. Which particular action will maximize happiness more than any other, in a particular situation, can be determined only by empirical investigation. Thus, it is possible that killing is always wrong—at least the utilitarian could consistently say so and thus be a pacifist; but *if* killing is always wrong, it is wrong not because killing is wrong per se but because it always and without exception leads to worse consequences than any other actions that could have been performed instead. Then the pacifist, if he is a consistent utilitarian, would have to go on to show in each instance that each and every act of killing is worse (leads to worse consequences) than any act of refraining from doing so—even when the man is a trigger-happy gunman who will kill dozens of people in a crowded street if he is not killed first.

That killing is worse in every instance would be extremely difficult—most people would say impossible—to prove.

Consider the syllogism:

The action which promotes the maximum happiness is right.

This action is the one which promotes the maximum happiness.

Therefore, *This action is right.*

The utilitarian gives undeviating assent only to the *first* of these three statements (the major premise); this statement is the chief article of his utilitarian creed, and he cannot abandon it without being inconsistent with his own doctrine. But this first premise is not enough to yield the third statement, which is the conclusion of the argument. To know that the conclusion is true, even granting that the major premise is, one must also know whether the second statement (the minor premise) is true; and the second statement is an empirical one, which cannot be verified by the philosopher sitting in his study but only by a thorough investigation of the empirical facts of the situation. Many people would accept the major premise (and thus be utilitarians) and yet disagree among themselves on the conclusion because they would disagree on the minor premise. They would agree that an act is right if it leads to maximum happiness, but they would not agree on whether this action or that one is the one which *will* in fact lead to the most happiness. They disagree about the empirical facts of the case, not in their utilitarian ethics. The disagreement could be resolved if both parties had a complete grasp of all the relevant empirical facts, for then they would know *which* action *would* lead to the most happiness. In many situations, of course, such agreement will never be reached because the consequences of people's actions (especially when they affect thousands of other people over a long period of time, as happens when war is declared) are so numerous and so complex that nobody will ever know them all. Such a disagreement will not be the fault of ethics, or of philosophy in general, but of the empirical world for being so complicated and subtle in its workings that the full consequences of our actions often can not be determined. Frequently it would take an omniscient deity to know which action in a particular situation was right. Finite human beings have to be content with basing their actions on estimates of probability.

According to utilitarianism, then, the traditional moral rules are justified for the most part because following them will lead to the best consequences far more often than violating them will; and that is why they are useful rules of thumb in human action. But, for the utilitarian, this is *all* they are—rules of thumb. They should never be used blindly, as a pat formula or inviolable rule subject to no exceptions, without an eye to the detailed consequences in each particular situation. The judge who condemned a man to

death in the electric chair for stealing \$1.95 (as in the case in Alabama in 1959) was probably not contributing to human happiness by inflicting this extreme penalty, even though he acted in accordance with the law of that state. The utilitarian would say that if a starving man steals a loaf of bread, as in Victor Hugo's *Les Misérables*, he should not be condemned for violating the rule "Do not steal"; in fact he probably did nothing morally wrong by stealing in this instance because the effects of not stealing would . . . have meant starvation and preserving a life (the utilitarian would say) is more important to human happiness than refraining from stealing a loaf of bread—especially since the man stole from one who was far from starving himself (the "victim" would never have missed it). He is probably blameless furthermore because the whole episode was made possible in the first place by a system of laws and a social structure which, by any utilitarian standard, were vicious in the extreme. (But see the effects of lawbreaking, below.)

Moral rules are especially useful when we have to act at once without being able adequately to weigh the consequences; for *usually* (as experience shows) better—i.e., more-happiness-producing—consequences are obtained by following moral rules than by not following them. If there is a drowning person whom you could rescue, you should do so without further investigation; for if you stopped to investigate his record, he would already have drowned. True, he might turn out to be a Hitler, but unless we have such evidence, we have to go by the probability that the world is better off for his being alive than his being dead. Again, there may be situations in which telling a lie will have better affects than telling the truth. But since, on the whole, lying has bad effects, we have to have special evidence that this situation is different before we are justified in violating the rule. If we have no time to gather such evidence, we should act on what is most probable, namely that telling a lie in this situation will produce consequences less good than telling the truth.

The utilitarian attitude toward moral rules is more favorable than might first appear because of the hidden, or subtle, or not frequently thought of, consequences of actions which at first sight would seem to justify a violation of the rules. One might consider *all* the consequences of the action and not just the immediate ones or the ones that happen to be the most conspicuous. For example: the utilitarian would not hold that it is *always* wrong to break a law, unless, he had good grounds for saying that breaking the law *always* leads to worse consequences than observing it. But if the law is a bad law to begin with or even if it is a good law on the whole but observing the law in this particular case would be deleterious to human happiness, then the law should be broken in this case. You would be morally justified, for example, in breaking the speed law in order to rush a

badly wounded person to a hospital. But in many situations (probably in most) in which the utilitarian criterion at first *seems* to justify the violation of a law, it does not really do so after careful consideration because of the far-flung consequences. For example, in a more typical instance of breaking the speed law, you might argue as follows: "It would make me happier if I were not arrested for the violation, and it wouldn't make the arresting officer any the less happy, in fact it would save him the trouble of writing out the ticket, so—why not? By letting me go, wouldn't the arresting officer be increasing the total happiness of the world by just a little bit, both his and mine, whereas by giving me a ticket he might actually decrease the world's happiness slightly?"

But happiness would be slightly increased only if one considers only the immediate situation. For one thing, by breaking the speed limit you are endangering the lives of others—you are less able to stop or to swerve out of the way in an emergency. Also those who see you speeding and escaping the penalty may decide to do the same thing themselves; even though you don't cause any accidents by your violation, *they* may do so after taking their cue from you. Moreover, lawbreaking may reduce respect for law itself; although there may well be unjust laws and many laws could be improved, it is usually better (has better consequences) to work for their repeal than to break them while they are still in effect. Every violation decreases the effectiveness of law, and we are surely better off having law than not having it at all—even the man who violently objects to a law and complains bitterly when he's arrested will invoke the law to protect *himself* against the violations of others. In spite of these cautions, utilitarianism does not say that one should *never* break a law but only that the consequences of doing so are far more often bad than good; a closer look at the consequences will show how true their reasoning is.

Notes

1. This term is most precisely defined by G. E. Moore in chapter 1, "Utilitarianism," of his book *Ethics*. [New York: Henry Holt and Company, 1912]. For the clearest and most rigorous statement of utilitarianism in its hedonistic form, see chapters 1 and 2 of [Moore's] book.
2. J. S. Mill, *Utilitarianism*. [ed. Oskar Piest (New York: Bobbs-Merrill, 1957; originally published 1863.), chap. 2.

2

The Best Action Is the One in Accord with Universal Rules

An important competitor to the utilitarian moral theory is the theory developed by Immanuel Kant. The fundamental ethical principle of Kantian theory is this: People should always be treated as ends, never as simply a means. In other words, it is wrong to ignore another person's legitimate desires and to use him or her just to get what you want. In the following essay James Rachels explains the Kantian view.

An example of Kantian thinking can be found in Stallman's argument that programmers should treat other computer users as persons whose desires should be respected, rather than as economic units (see "The GNU Manifesto").

James Rachels

The great German philosopher Immanuel Kant thought that human beings occupy a special place in creation. Of course he was not alone in thinking this. It is an old idea: from ancient times, humans have considered themselves to be essentially different from all other creatures—and not just different but *better*. In fact, humans have traditionally thought themselves to be quite fabulous. Kant certainly did. [I]n his view, human beings have "an intrinsic worth, i.e., *dignity*," which makes them valuable "above all price." Other animals, by contrast, have value only insofar as they serve human purposes. In his *Lecture on Ethics* (1779), Kant said:

But so far as animals are concerned, we have no direct duties. Animals . . . are there merely as means to an end. That end is man.

We can, therefore, use animals in any way we please. We do not even have a "direct duty" to refrain from torturing them. Kant admits that it is probably wrong to torture them, but the reason is not that *they* would be hurt; the reason is only that we might suffer indirectly as a result of it, because "he

From James Rachels, *The Elements of Moral Philosophy* © 1986. Reprinted by permission of The McGraw-Hill Companies.

who is cruel to animals becomes hard also in his dealings with men." Thus [i]n Kant's view, mere animals have no importance at all. Human beings are, however, another story entirely. According to Kant, humans may never be "used" as means to an end. He even went so far to suggest that this is the ultimate law of morality.

Like many other philosophers, Kant believed that morality can be summed up in one ultimate principle, from which all our duties and obligations are derived. He called this principle *The Categorical Imperative*. In the *Groundwork of the Metaphysics of Morals* (1785) he expressed it like this:

Act only according to that maxim by which you can at the same time will that it should become a universal law.

However, Kant also gave *another* formulation of The Categorical Imperative. Later in the same book, he said that the ultimate moral principle may be understood as saying:

Act so that you treat humanity, whether in your own person or in that of another, always as an end and never as a means only.

Scholars have wondered ever since why Kant thought these two rules were equivalent. They *seem* to express very different moral conceptions. Are they, as he apparently believed, two versions of the same basic idea, or are they really different ideas? We will not pause over this question. Instead we will concentrate here on Kant's belief that morality requires us to treat persons "always as an end and never as a means only." What exactly does this mean, and why did he think it true?

When Kant said that the value of human beings "is above all price," he did not intend this as mere rhetoric but as an objective judgment about the place of human beings in the scheme of things. There are two important facts about people that, in his view, support his judgment.

First, because people have desires and goals, other things have value *for them*, in relation to *their* projects. Mere "things" (and this includes nonhuman animals, whom Kant considered unable to have self-conscious desires and goals) have value only as means to ends, and it is human ends that *give* them value. Thus if you want to become a better chess player, a book of chess instruction will have value for you; but apart from such ends the book has no value. Or if you want to travel about, a car will have value for you; but apart from this desire the car will have no value.

Second, and even more important, humans have "an intrinsic worth, i.e., *dignity*," because they are *rational agents*—that is, free agents capable of making their own decisions, setting their own goals, and *guiding their conduct by reason*. Because the moral law is the law of reason, rational beings are the embodiment of the moral law itself. The only way that moral good-

ness can exist at all in the world is for rational creatures to apprehend what they should do and, acting from a sense of duty, do it. This, Kant thought, is the *only* thing that has “moral worth.” Thus if there were no rational beings, the moral dimension of the world would simply disappear.

It makes no sense, therefore, to regard rational beings merely as one kind of valuable thing among others. They are the beings *from whom* mere “things” have value, and they are the beings whose conscientious actions have moral worth. So Kant concludes that their value must be absolute, and not comparable to the value of anything else.

If their value is “beyond all price,” it follows that rational beings must be treated “always as an end, and never as a means only.” This means, on the most superficial level, that we have a strict duty of beneficence toward other persons: we must strive to promote their welfare; we must respect their rights, avoid harming them, and generally “endeavor, so far as we can, to further the ends of others.”

But Kant’s idea also has a somewhat deeper implication. The beings we are talking about are *rational* beings, and “treating them as ends-in-themselves” means *respecting their rationality*. Thus we may never *manipulate* people, or *use* people, to achieve our purposes, no matter how good those purposes may be. Kant gives this example, which is similar to an example he uses to illustrate the first version of his categorical imperative. Suppose you need money, and so you want a “loan,” but you know you will not be able to repay it. In desperation, you consider making a false promise (to repay) in order to trick a friend into giving you the money. May you do this? Perhaps you need the money for a good purpose—so good, in fact, that you might convince yourself the lie would be justified. Nevertheless, if you lied to your friend, you would merely be manipulating him and using him “as means.”

On the other hand, what would it be like to treat your friend “as an end”? Suppose you told the truth, that you need the money for a certain purpose but will not be able to repay it. Then your friend could make up his own mind about whether to let you have it. He could exercise his own powers of reason, consulting his own value and wishes, and make a free, autonomous choice. If he did decide to give the money for this purpose, he would be choosing to make that purpose *his own*. Thus you would not merely be using him as a means to achieving *your* goal. This is what Kant meant when he said, “Rational beings . . . must always be esteemed at the same time as ends, i.e., only as beings who must be able to contain in themselves the end of the very same action.”

Now Kant’s conception of human dignity is not easy to grasp; it is, in fact, probably the most difficult notion discussed [here]. We need to find a way to make the idea clearer. In order to do that, we will consider in some detail one of its most important applications—this may be better than a dry,

theoretical discussion. Kant believed that if we take the idea of human dignity seriously, we will be able to understand the practice of criminal punishment in a new and revealing way.

.

On the face of it, it seems unlikely that we could describe punishing someone as “respecting him as a person” or as “treating him as an end-in-himself.” How could taking away someone’s freedom, by sending him to prison, be a way of “respecting” him? Yet that is exactly what Kant suggests. Even more paradoxically, *he implies that executing someone may also be a way of treating him “as an end.”* How can this be?

Remember that, for Kant, treating someone as an “end-in-himself” means treating him *as a rational being*. Thus we have to ask, What does it mean to treat someone as a rational being? Now a rational being is someone who is capable of reasoning about his conduct and who freely decides what he will do, on the basis of his own rational conception of what is best. Because he has these capacities, a rational being is *responsible* for his actions. We need to bear in mind the difference between:

1. *Treating someone as a responsible being*

and

2. *Treating someone as a being who is not responsible for his conduct.*

Mere animals, who lack reason, are not responsible for their actions; nor are people who are mentally “sick” and not in control of themselves. In such cases it would be absurd to try to “hold them accountable.” We could not properly feel gratitude or resentment toward them, for they are not responsible for any good or ill they cause. Moreover, we cannot expect them to understand why we treat them as we do, any more than they understand why they behave as they do. So we have no choice but to deal with them by manipulating them, rather than by addressing them as autonomous individuals. When we spank a dog who has urinated on the rug, for example, we may do so in an attempt to prevent him from doing it again—but we are merely trying to “train” him. We could not reason with him even if we wanted to. The same goes for mentally “sick” humans.

On the other hand, rational beings are responsible for their behavior and so may properly be “held accountable” for what they do. We may feel gratitude when they behave well, and resentment when they behave badly. Reward and punishment—not “training” or other manipulation—are the natural expression of this gratitude and resentment. Thus in punishing people, we are *holding them responsible* for their actions, in a way in which we cannot hold mere animals responsible. We are responding to them not as —

people who are "sick" or who have no control over themselves, but as people who have freely chosen their evil deeds.

Furthermore, in dealing with responsible agents, we may properly allow *their conduct* to determine, at least in part, how we respond to them. If someone has been kind to you, you may respond by being generous in return; and if someone is nasty to you, you may also take that into account in deciding how to deal with him or her. And why shouldn't you? Why should you treat everyone alike, regardless of how *they* have chosen to behave?

Kant gives this last point a distinctive twist. There is [i]n his view, a deep logical reason for responding to other people "in kind." The first formulation of The Categorical Imperative comes into play here. When we decide what to do, we in effect proclaim our wish that our conduct be made into a "universal law." Therefore, when a rational being decides to treat people in a certain way, he decrees that in his judgment *this is the way people are to be treated*. Thus if we treat him the same way in return, we are doing nothing more than treating him as *he has decided* people are to be treated. If he treats others badly, and we treat him badly, we are complying with his own decision. (Of course, if he treats others well, and we treat him well in return, we are also complying with the choice he has made.) We are allowing *him* to decide how he is to be treated—and so we are, in a perfectly clear sense, respecting his judgment, by allowing it to control our treatment of him. Thus Kant says of the criminal, "His own evil deed draws the punishment upon himself."

3

The Best Action Is the One That Exercises the Mind's Faculties

The question of what constitutes ethical action has been a subject of discussion for millennia, as this selection from Aristotle (4th c. BCE) shows. Aristotle's answer is that ethical action consists in the active exercise of the mind's faculties. Those who are most involved with computerization, for whom reasoning and intellectual challenge are

From Aristotle, *Ethics for English Readers*, trans. H. Rackham. Oxford: Basil Blackwell, Publisher, 1952. Reprinted by permission.

highly valued, may well agree with Aristotle. (See also Dreyfus and Dreyfus, "Using Computers as Means, Not Ends.")

Aristotle

... Inasmuch as all study and all deliberate action is aimed at some good object, **let us state what is the good which is in our view the aim of political science**, and what is the highest of the goods obtainable by action.

Now as far as the name goes there is virtual agreement about this among the vast majority of mankind. Both ordinary people and persons of trained mind define the good as happiness. But as to what constitutes happiness opinions differ: **the answer given by ordinary people is not the same as the verdict of the philosopher**. Ordinary men identify happiness with something obvious and visible, such as pleasure or wealth or honor—everybody gives a different definition, and sometimes the same person's own definition alters: when a man has fallen ill he thinks that happiness is health, if he is poor he thinks it is wealth. And when people realize their own ignorance they regard with admiration those who propound some grand theory that is above their heads. The view has been held by some thinkers¹ that besides the many good things alluded to above there also exists something that is good in itself, which is the fundamental cause of the goodness of all the others.

Now to review the whole of these opinions would perhaps be a rather thankless task. It may be enough to examine those that are most widely held, or that appear to have some considerable argument in their favor. . . .

Reasons for doubting whether enjoyment, fame, virtue, or wealth is the whole good

To judge by men's mode of living, the mass of mankind think that good and happiness consist in pleasure, and consequently are content with a life of mere enjoyment. **There are in fact three principal modes of life—the one just mentioned, the life of active citizenship, and the life of contemplation**. The masses, being utterly servile, obviously prefer the life of mere cattle; and indeed they have some reason for this, inasmuch as many men of high station share the tastes of Sardanapalus.² The better people, on the other hand, and men of action, give the highest value to honor, since honor may be said to be the object aimed at in a public career. Nevertheless, it would seem that honor is a more superficial thing than the good which we are in search of, because honor seems to depend more on the people who render it than on the person who receives it, whereas we dimly feel that good must be something inherent in oneself and inalienable. Moreover, men's object

in pursuing honor appears to be to convince themselves of their own worth; at all events they seek to be honored by persons of insight and by people who are well acquainted with them, and to be honored for their merit. It therefore seems that at all events in the opinions of these men goodness is more valuable than honor, and probably one may suppose that it has a better claim than honor to be deemed the end at which the life of politics aims. But even virtue appears to lack completeness as an end, inasmuch as it seems to be possible to possess it while one is asleep or living a life of perpetual inactivity, and moreover one can be virtuous and yet suffer extreme sorrow and misfortune; but nobody except for the sake of maintaining a paradox would call a man happy in those circumstances.

However, enough has been said on this topic, which has indeed been sufficiently discussed in popular treatises.

The third life is the life of contemplation, which we shall consider later.

The life of money-making is a cramped way of living, and clearly wealth is not the good we are in search of, as it is only valuable as a means to something else. Consequently a stronger case might be made for the objects previously specified, because they are valued for their own sake; but even they appear to be inadequate, although a great deal of discussion has been devoted to them. . . .

Reaffirmation that the good is the ultimate and self-sufficient object of desire and that "happiness" is the good

What then is the precise nature of the practicable good which we are investigating? It appears to be one thing in one occupation or profession and another in another: the object pursued in medicine is different from that of military science, and similarly in regard to the other activities. What definition of the term "good" then is applicable to all of them? Perhaps "the object for the sake of attaining which all the subsidiary activities are undertaken." The object pursued in the practice of medicine is health, in a military career victory, in architecture a building—one thing in one pursuit and another in another, but in every occupation and every pursuit it is the end aimed at, since it is for the sake of this that the subsidiary activities in all these pursuits are undertaken. Consequently if there is some one thing which is the end and aim of all practical activities whatsoever, that thing, or if there are several, those things, will constitute the practicable good.

Our argument has therefore come round again by a different route to the point reached before. We must endeavor to render it yet clearer.

Now the objects at which our actions aim are manifestly several, and some of these objects, for instance money, and instruments in general, we adopt as means to the attainment of something else. This shows that not all the objects we pursue are final ends. But the greatest good manifestly is a final end. Consequently if there is only one thing which is final, that will be

the object for which we are now seeking, or if there are several, it will be that one among them which possesses the most complete finality.

Now a thing that is pursued for its own sake we pronounce to be more final than one pursued as a means to some other thing, and a thing that is never desired for the sake of something else we call more final than those which are desired for the sake of something else as well as for their own sake. In fact the absolutely final is something that is always desired on its own account and never as a means for obtaining something else. Now this description appears to apply in the highest degree to happiness, since we always desire happiness for its own sake and never on account of something else; whereas honor and pleasure and intelligence and each of the virtues, though we do indeed desire them on their own account as well, for we should desire each of them even if it produced no external result, we also desire for the sake of happiness, because we believe that they will bring it to us, whereas nobody desires happiness for the sake of those things, not for anything else but itself.

The same result seems to follow from a consideration of the subject of self-sufficiency, which is felt to be a necessary attribute of the final good. The term self-sufficient denotes not merely being sufficient for oneself alone, as if one lived the life of a hermit, but also being sufficient for the needs of one's parents and children and wife, and one's friends and fellow-countrymen in general, inasmuch as man is by nature a social being.

Yet we are bound to assume some limit in these relationships, since if one extends the connection to include one's children's children and friends' friends, it will go on *ad infinitum*. But that is a matter which must be deferred for later consideration. Let us define self-sufficiency as the quality which makes life to be desirable and lacking in nothing even when considered by itself; and this quality we assume to belong to happiness. Moreover, when we pronounce happiness to be the most desirable of all things, we do not mean that it stands as one in a list of good things—were it so, it would obviously be more desirable in combination with even the smallest of the other goods, inasmuch as that addition would increase the total of good, and of two good things the larger must always be the more desirable.

Thus it appears that happiness is something final and complete in itself, as being the aim and end of all practical activities whatever.

For a more specific conception of the kind of "happiness" which is the good, we do well to examine whether nature intended man for anything, as it intended the eye for sight. What is distinctive of man is reason, so the happiness which is the good must be the exercise of reason in living.

Possibly, however, the student may feel that the statement that happiness is the greatest good is a mere truism, and he may want a clearer explanation

of what the precise nature of happiness is. This may perhaps be achieved by ascertaining what is the proper function of man. In the case of flute players or sculptors or other artists, and generally of all persons who have a particular work to perform, it is felt that their good and their well-being are found in that work. It may be supposed that this similarly holds good in the case of a human being, if we may assume that there is some work which constitutes the proper function of a human being as such. Can it then be the case that whereas a carpenter and a shoemaker have definite functions or businesses to perform, a man as such has none, and is not designed by nature to perform any function? Should we not rather assume that, just as the eye and hand and foot and every part of the body manifestly have functions assigned to them, so also there is a function that belongs to a man, over and above all the special functions that belong to his members? If so, what precisely will that function be? It is clear that the mere activity of living is shared by man even with the vegetable kingdom, whereas we are looking for some function that belongs specially to man. We must therefore set aside the vital activity of nutrition and growth. Next perhaps comes the life of the senses; but this also is manifestly shared by the horse and the ox and all the animals. There remains therefore what may be designated the practical life of the rational faculty.

But the term "rational" life has two meanings: it denotes both the mere possession of reason, and its active exercise. Let us take it that we here mean the latter, as that appears to be the more proper signification of the term. Granted then that the special function of man is the active exercise of the mind's faculties in accordance with rational principle, or at all events not in detachment from rational principle, and that the function of anything, for example, a harper, is generally the same as the function of a good specimen of that thing, for example a good harper (the specification of the function merely being augmented in the latter case with the statement of excellence—a harper is a man who plays the harp, a good harper one who plays the harp well)—granted, I say, the truth of these assumptions, it follows that the good of man consists in the active exercise of the faculties in conformity with excellence or virtue, or if there are several virtues, in conformity with the best and most perfect among them.

Notes

1. Plato and the Academy.
2. A mythical Assyrian king; two versions of his epitaph are recorded, one containing the words "Eat, drink, play, since all else is not worth that snap of the fingers," the other ending "I have what I ate, and the delightful deeds of wantonness and love in which I shared; but all my wealth is vanished."

Professional Ethics

How can computing professionals apply ethical decision making to their work? Because of the pervasiveness and power of information technology, many ethical dilemmas emerge in the field of computing.

This section examines the ways in which these dilemmas can be negotiated. For further study of computing ethics, you may also visit the Brookings Computer Ethics Institute website, www.brook.edu/its/cei/cei_hp.htm.